

Linear Regression

Dit-Yan Yeung

Department of Computer Science and Engineering
Hong Kong University of Science and Technology

COMP 4211: Machine Learning (Fall 2022)

- 1 Introduction
- 2 Linear Regression
- 3 Extensions of Linear Regression
- 4 Regularization
- 5 Performance Metrics

Regression Revisited

- Given a training set $\mathcal{S} = \{(\mathbf{x}^{(\ell)}, \mathbf{y}^{(\ell)})\}_{\ell=1}^N$ of N **labeled examples** each of which is in the form of an input-output pair.
- The problem is to estimate the parameters \mathbf{w} in a **regression function** $f(\mathbf{x}; \mathbf{w})$ using \mathcal{S} such that the **predicted output** $f(\mathbf{x}^{(\ell)}; \mathbf{w})$ for each input $\mathbf{x}^{(\ell)}$ is (usually) close to the **actual output** $\mathbf{y}^{(\ell)}$. Moreover, we want this to hold also for unseen examples sampled from the same data distribution.
- When the output \mathbf{y} (with the superscript ℓ dropped for notational simplicity) is multivariate (i.e., \mathbf{y} is a vector), it is a **multi-output regression** problem.
- A more common form, which will be our focus here, is when \mathbf{y} is univariate (i.e., \mathbf{y} degenerates to a scalar). We denote the output by y instead.
- The input $\mathbf{x} = (x_1, \dots, x_d)^\top$ is d -dimensional (usually $d > 1$).

Linear Regression Function

- In **linear regression**, the regression function is simply a linear function:

$$f(\mathbf{x}; \mathbf{w}) = w_0 + w_1x_1 + \cdots + w_dx_d = \mathbf{w}^\top \tilde{\mathbf{x}} = \tilde{\mathbf{x}}^\top \mathbf{w},$$

where $\mathbf{w} = (w_0, \dots, w_d)^\top$ denotes the **parameter vector** and $\tilde{\mathbf{x}} = (x_0 = 1, x_1, \dots, x_d)^\top$ is the **augmented input** with the constant 1 introduced as an additional dimension x_0 .

- The weight w_0 is called the **bias** term which serves as an offset.
- The **learning** problem is to find the best \mathbf{w} according to some performance measure using the training set \mathcal{S} .

Squared Loss

- Like many (though not all) machine learning models, a common way to learn the parameters \mathbf{w} of the linear regression function $f(\mathbf{x}; \mathbf{w})$ is to define a **loss function** $L(\mathbf{w}; \mathcal{S})$ as a function of \mathbf{w} and then minimize $L(\mathbf{w}; \mathcal{S})$ w.r.t. \mathbf{w} .
- The most common loss function for regression problems is the **squared loss** (a.k.a. **quadratic loss**):

$$\begin{aligned} L(\mathbf{w}; \mathcal{S}) &= \sum_{\ell=1}^N \left(f(\mathbf{x}^{(\ell)}; \mathbf{w}) - y^{(\ell)} \right)^2 \\ &= \sum_{\ell=1}^N \left(w_0 + w_1 x_1^{(\ell)} + \cdots + w_d x_d^{(\ell)} - y^{(\ell)} \right)^2. \end{aligned}$$

A Special Case ($d = 1$) for Illustration

- Squared loss:

$$L(\mathbf{w}; \mathcal{S}) = \sum_{\ell=1}^N \left(w_0 + w_1 x_1^{(\ell)} - y^{(\ell)} \right)^2.$$

- Since $L(\mathbf{w}; \mathcal{S})$ is quadratic in \mathbf{w} , the **optimal solution** $\hat{\mathbf{w}}$ that minimizes $L(\mathbf{w}; \mathcal{S})$ is **unique** and can be found in **closed form** using the method of **least squares**.
- Setting the derivatives of $L(\mathbf{w}; \mathcal{S})$ w.r.t. w_0 and w_1 to 0 gives two **linear equations**:

$$\begin{aligned} Nw_0 + w_1 \sum_{\ell=1}^N x_1^{(\ell)} &= \sum_{\ell=1}^N y^{(\ell)} \\ w_0 \sum_{\ell=1}^N x_1^{(\ell)} + w_1 \sum_{\ell=1}^N \left(x_1^{(\ell)} \right)^2 &= \sum_{\ell=1}^N x_1^{(\ell)} y^{(\ell)}. \end{aligned}$$

Linear System in Matrix Form

- The linear equations can be expressed as:

$$\mathbf{A}\mathbf{w} = \begin{bmatrix} N & \sum_{\ell} x_1^{(\ell)} \\ \sum_{\ell} x_1^{(\ell)} & \sum_{\ell} (x_1^{(\ell)})^2 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} = \begin{bmatrix} \sum_{\ell} y^{(\ell)} \\ \sum_{\ell} x_1^{(\ell)} y^{(\ell)} \end{bmatrix} = \mathbf{b}.$$

- Least squares estimate expressed in closed form:

$$\hat{\mathbf{w}} = \mathbf{A}^{-1}\mathbf{b},$$

assuming that \mathbf{A} is invertible.

General Case ($d \geq 1$)

- Let the input and output parts of the N examples be expressed as:

$$\mathbf{X} = \begin{bmatrix} 1 & x_1^{(1)} & \cdots & x_d^{(1)} \\ 1 & x_1^{(2)} & \cdots & x_d^{(2)} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_1^{(N)} & \cdots & x_d^{(N)} \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(N)} \end{bmatrix}.$$

- The method of least squares gives a set of $d + 1$ linear equations which can be expressed in matrix form as:

$$\mathbf{A}\mathbf{w} = (\mathbf{X}^\top \mathbf{X}) \mathbf{w} = \mathbf{X}^\top \mathbf{y} = \mathbf{b}.$$

- Least squares estimate:

$$\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y},$$

assuming that $\mathbf{X}^\top \mathbf{X}$ is invertible.

Alternative Derivation using Multivariable Calculus

- The squared loss can also be expressed as:

$$\begin{aligned}
 L(\mathbf{w}; \mathcal{S}) &= \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 \\
 &= (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) \\
 &= \mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{y}^\top \mathbf{X} \mathbf{w} + \mathbf{y}^\top \mathbf{y},
 \end{aligned}$$

where $\|\mathbf{v}\| = (\sum_i v_i^2)^{1/2}$ denotes the **L_2 norm** of a vector \mathbf{v} .

- To minimize $L(\mathbf{w}; \mathcal{S})$, we apply multivariate calculus to differentiate $L(\mathbf{w}; \mathcal{S})$ w.r.t. \mathbf{w} and set the derivative to the zero vector $\mathbf{0}$ to get:

$$\begin{aligned}
 2\mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{X}^\top \mathbf{y} &= \mathbf{0} \\
 (\mathbf{X}^\top \mathbf{X}) \mathbf{w} &= \mathbf{X}^\top \mathbf{y},
 \end{aligned}$$

which is the same as the matrix form of the system of $d + 1$ linear equations obtained above.

Complexity Considerations

- The closed-form solution requires inverting $\mathbf{X}^\top \mathbf{X}$ which is a $(d + 1) \times (d + 1)$ matrix.
- When d is large, instead of resorting to a closed-form solution, an alternative approach is to estimate $\hat{\mathbf{w}}$ **iteratively** like the **logistic regression** model for classification (to be discussed in the next topic).

Nonlinear Extensions

- For solving more complicated problems, **nonlinear** regression functions are needed.
- Different approaches for nonlinear extension:
 - ① **Explicitly** adding more input dimensions (which depend nonlinearly on the original input dimensions) and applying linear regression to the **expanded input**
 - ② Applying a **nonlinear regression function** to the original input
 - ③ **Implicitly** transforming the original input nonlinearly to a new space and applying a linear model to the **transformed input**
- We consider the first approach here and leave the other two for some later topics.

Polynomial Regression

- One common approach is to introduce **higher-order terms** as additional input dimensions, e.g., x_i^2 , $x_i x_j$, $x_i x_j^2 x_k$.
- For notational simplicity, we only consider the $d = 1$ case and write the (only) input dimension x_1 simply as x .
- Polynomial function of **degree** d :

$$f(\mathbf{x}; \mathbf{w}) = w_0 + w_1 x + \cdots + w_d x^d = \mathbf{w}^\top \tilde{\mathbf{x}} = \tilde{\mathbf{x}}^\top \mathbf{w},$$

where $\tilde{\mathbf{x}} = (1, x, \dots, x^d)^\top$ is like the augmented input used for linear regression.

Other Additional Input Dimensions

- Besides adding higher-order terms, more general transformations of the original input dimensions may also be introduced as additional input dimensions.
- Very often, **feature engineering** that uses domain knowledge to define **application-specific features** is applied.
- The method of least squares for linear regression can also be used here (for both polynomial regression and more general extensions) to obtain a **closed-form solution**.
- For the subsequent discussions, we use the same generic formulation regardless of whether additional input dimensions are introduced.

Model Overfitting

- If the training set $\mathcal{S} = \{(\mathbf{x}^{(\ell)}, y^{(\ell)})\}_{\ell=1}^N$ is small compared to the number of parameters in the linear regression function $f(\mathbf{x}; \mathbf{w}) = w_0 + w_1x_1 + \cdots + w_dx_d$, **overfitting** may occur.
- When overfitting occurs, it is common to find large magnitudes (i.e., absolute values) in at least some of the parameters.
- One common solution to the overfitting problem is to prevent the parameters from growing excessively large in magnitude.

Regularization

- **Regularization** is an approach which modifies the original loss function by adding one or more penalty terms, called **regularizers**, that penalize large parameter magnitudes.
- **Regularized loss function** based on L_2 regularization (a.k.a. **Tikhonov regularization**):

$$\begin{aligned} L_\lambda(\mathbf{w}; \mathcal{S}) &= L(\mathbf{w}; \mathcal{S}) + \lambda \|\mathbf{w}\|^2 \\ &= \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 + \lambda \|\mathbf{w}\|^2 \\ &= \mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - 2 \mathbf{y}^\top \mathbf{X} \mathbf{w} + \mathbf{y}^\top \mathbf{y} + \lambda \mathbf{w}^\top \mathbf{w}, \end{aligned}$$

where $\lambda > 0$ is called the **regularization parameter**.

- In practice, not regularizing w_0 usually gives better result because w_0 serves as an **offset** which is not multiplied with any input variable.

Closed-Form Solution with L_2 Regularization

- By differentiating $L_\lambda(\mathbf{w}; \mathcal{S})$ w.r.t. \mathbf{w} and setting the derivative to the zero vector $\mathbf{0}$, we obtain:

$$\begin{aligned} 2\mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{X}^\top \mathbf{y} + 2\lambda \mathbf{w} &= \mathbf{0} \\ (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}) \mathbf{w} &= \mathbf{X}^\top \mathbf{y}, \end{aligned}$$

where \mathbf{I} is the identity matrix.

- The least squares estimate can also be obtained in closed form:

$$\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y}.$$

- Note that $\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}$ is always invertible for any $\lambda > 0$.
- Linear regression with L_2 regularization (a.k.a. **ridge regression**) degenerates to the ordinary linear regression (without regularization) when $\lambda = 0$.


How to Set λ ?

- Though not the only method, **cross validation** (or, more correctly, called **holdout validation**) is commonly used by training a model (using a regularized loss function with a specific value of λ) on a **training set** and validating the trained model on a separate **validation set** (which mimics the **test set** not available during model training).
- Different values of λ are used to obtain different least squares estimates. The value that gives the smallest error rate in the validation set is chosen.

Overfitting and Underfitting

- Suppose the training set is small compared to the number of model parameters.
- **Overfitting:**
 - May occur if λ is too small
 - Small training error
 - Large validation error
- **Underfitting:**
 - May occur if λ is too large
 - Large training error
 - Large validation error

Typical Training and Validation Error Curves



Other Regularizers

- Many other regularizers can also be defined.
- For example, instead of using the L_2 norm, the L_p norm for some other value of p (e.g., 1 or 0) has also been used:

$$\|\mathbf{v}\|_p = \left(\sum_i |v_i|^p \right)^{1/p}.$$

(The L_2 norm may also be written explicitly as $\|\cdot\|_2$.)

- Linear regression with L_1 regularization, also called **LASSO** (least absolute shrinkage and selection operator), favors **sparse** solutions with all but a small number of dimensions equal to 0.
- Although the L_1 norm is also **convex** like the L_2 norm, there is no closed-form solution for LASSO. **Iterative algorithms** are needed for estimating the parameters.

Mean Squared Error

- A common performance metric for regression problems is the **mean squared error (MSE)**:

$$MSE = \frac{1}{N} \sum_{\ell=1}^N \left(f(\mathbf{x}^{(\ell)}; \mathbf{w}) - y^{(\ell)} \right)^2,$$

which is similar to the squared loss but with two differences:

- MSE can be used for the **validation set** and **test set** in addition to the **training set**.
- MSE measures the **mean** over all the examples in the set, not the **sum**.
- Instead of using MSE, it is more common to use the **root mean squared error (RMSE)**, which is the square root of MSE, to bring it back to the same level of prediction error for easier interpretation.

R^2 Score

- Another commonly used performance metric for regression problems is the **coefficient of determination** or R^2 score:

$$R^2 = 1 - \frac{\sum_{\ell=1}^N (f(\mathbf{x}^{(\ell)}; \mathbf{w}) - y^{(\ell)})^2}{\sum_{\ell=1}^N (\bar{y} - y^{(\ell)})^2},$$

where

$$\bar{y} = \frac{1}{N} \sum_{\ell=1}^N y^{(\ell)}.$$

- The best possible R^2 score is 1 when the corresponding MSE is 0.
- When the model always predicts the mean value of y , the R^2 score will be equal to 0.
- Negative values are also possible because the model can have arbitrarily large MSE.