

# Support Vector Machines


Dit-Yan Yeung

Department of Computer Science and Engineering  
Hong Kong University of Science and Technology


COMP 4211: Machine Learning (Fall 2022)

- 1 Introduction
- 2 Hard-Margin Support Vector Machine
- 3 Soft-Margin Support Vector Machine
- 4 Alternative Formulation
- 5 Kernel Extension
- 6 Further Study


# Given a Data Set ...



## ... Which Separating Hyperplane is the Best?




# Optimal Separating Hyperplane



- **Margin** of a separating hyperplane: **distance** to the separating hyperplane from the data point **closest** to it.
- Relationship between **margin** and **generalization**: There exist theoretical results showing that the separating hyperplane with the **largest margin** generalizes best (i.e., has **smallest generalization error**).

# Canonical Optimal Separating Hyperplane



- **Hard-margin case**: data points from the two classes are assumed to be **linearly separable**.
- With proper scaling of  $\mathbf{w}$  and  $w_0$ , the points closest to the hyperplane satisfy  $|\mathbf{w}^T \mathbf{x} + w_0| = 1$ . Such a hyperplane is called a **canonical separating hyperplane**.
- The one that **maximizes the margin** is called the **canonical optimal separating hyperplane**.

## Canonical Optimal Separating Hyperplane (2)

- Let  $\mathbf{x}^{(1)}$  and  $\mathbf{x}^{(2)}$  be two closest points, one on each side of the hyperplane.
- Note that  $\mathbf{w}$  is a **normal vector** to the hyperplane (i.e., its direction is perpendicular to that of the hyperplane) and

$$\mathbf{w}^\top \mathbf{x}^{(1)} + w_0 = +1$$

$$\mathbf{w}^\top \mathbf{x}^{(2)} + w_0 = -1,$$

which imply

$$\mathbf{w}^\top (\mathbf{x}^{(1)} - \mathbf{x}^{(2)}) = 2.$$

Hence the **margin** can be given by

$$\gamma = \frac{1}{2} \frac{\mathbf{w}^\top (\mathbf{x}^{(1)} - \mathbf{x}^{(2)})}{\|\mathbf{w}\|} = \frac{1}{\|\mathbf{w}\|}.$$

- Thus, **maximizing the margin** is equivalent to **minimizing  $\|\mathbf{w}\|$** .

# Inequality Constraints

- For all data points in the sample  $\mathcal{X} = \{(\mathbf{x}^{(\ell)}, y^{(\ell)})\}$ , we want  $\mathbf{w}$  and  $w_0$  to satisfy

$$\mathbf{w}^\top \mathbf{x}^{(\ell)} + w_0 \begin{cases} \geq +1 & \text{if } y^{(\ell)} = +1 \\ \leq -1 & \text{if } y^{(\ell)} = -1. \end{cases}$$

- Equivalent form of **inequality constraints**:

$$y^{(\ell)}(\mathbf{w}^\top \mathbf{x}^{(\ell)} + w_0) \geq 1. \quad (1)$$

- Instead of using inequality constraints

$$y^{(\ell)}(\mathbf{w}^\top \mathbf{x}^{(\ell)} + w_0) \geq 0,$$

which only require the data points to lie on the right side of the hyperplane, the constraints in (1) also want them to be some distance away for **better generalization**.



# Primal Optimization Problem

- Primal optimization problem:

$$\begin{array}{ll} \text{Minimize} & \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{subject to} & y^{(\ell)}(\mathbf{w}^\top \mathbf{x}^{(\ell)} + w_0) \geq 1, \forall \ell. \end{array}$$

- This is a **quadratic programming (QP)** problem, or a **quadratic program**, which is a type of **convex optimization** problem.
- In optimization theory, it is very common and sometimes advantageous to turn a primal problem into a **dual problem** and then solve the latter instead.
- In our case, it also turns out to be more convenient to solve the dual problem (whose complexity depends on the sample size  $N$ ) rather than the primal problem directly (whose complexity depends on the dimensionality  $d$ ). The dual problem also makes it easy for a **nonlinear** extension using **kernel functions**.

## Dual Optimization Problem

- By introducing  $N$  **Lagrange multipliers**,  $\{\alpha_\ell\}_{\ell=1}^N$ , one for each training data point, we can turn the primal problem into a dual problem (with details of the derivation skipped). The Lagrange multipliers are known as **dual variables** in the dual problem.
- Dual optimization problem:**

$$\begin{aligned} \text{Maximize} \quad & \sum_{\ell} \alpha_{\ell} - \frac{1}{2} \sum_{\ell} \sum_{\ell'} \alpha_{\ell} \alpha_{\ell'} y^{(\ell)} y^{(\ell')} (\mathbf{x}^{(\ell)})^{\top} \mathbf{x}^{(\ell')} \\ \text{subject to} \quad & \sum_{\ell} \alpha_{\ell} y^{(\ell)} = 0 \text{ and } \alpha_{\ell} \geq 0, \forall \ell. \end{aligned}$$

- This is also a **QP** problem, but its complexity depends on the sample size  $N$  (rather than the input dimensionality  $d$ ):
  - Time complexity:**  $O(N^3)$
  - Space complexity:**  $O(N^2)$

# Support Vectors

- At the optimal solution, most of the dual variables vanish with  $\alpha_\ell = 0$ . They are points lying beyond the margin with no effect on the hyperplane.
- **Support vectors**:  $\mathbf{x}^{(\ell)}$  with  $\alpha_\ell > 0$ , hence the name **support vector machine (SVM)**.

# Computation of Primal Variables from Dual Variables

- From the (skipped) derivation of the dual problem, we get

$$\mathbf{w} = \sum_{\ell=1}^N \alpha_{\ell} y^{(\ell)} \mathbf{x}^{(\ell)} = \sum_{\mathbf{x}^{(\ell)} \in \mathcal{SV}} \alpha_{\ell} y^{(\ell)} \mathbf{x}^{(\ell)},$$

where  $\mathcal{SV}$  denotes the set of support vectors.

- The support vectors must lie on the margin, so they should satisfy

$$y^{(\ell)}(\mathbf{w}^{\top} \mathbf{x}^{(\ell)} + w_0) = 1 \quad \text{or} \quad w_0 = y^{(\ell)} - \mathbf{w}^{\top} \mathbf{x}^{(\ell)}.$$

For numerical stability, all support vectors are used to compute  $w_0$ :

$$w_0 = \frac{1}{|\mathcal{SV}|} \sum_{\mathbf{x}^{(\ell)} \in \mathcal{SV}} (y^{(\ell)} - \mathbf{w}^{\top} \mathbf{x}^{(\ell)}).$$

# Discriminant Function

- Discriminant function:

$$\begin{aligned} g(\mathbf{x}) &= \mathbf{w}^\top \mathbf{x} + w_0 \\ &= \left( \sum_{\mathbf{x}^{(\ell)} \in \mathcal{SV}} \alpha_\ell y^{(\ell)} \mathbf{x}^{(\ell)} \right)^\top \mathbf{x} + \frac{1}{|\mathcal{SV}|} \sum_{\mathbf{x}^{(\ell)} \in \mathcal{SV}} (y^{(\ell)} - \mathbf{w}^\top \mathbf{x}^{(\ell)}). \end{aligned}$$

- Classification rule during testing:


$$\text{Choose } \begin{cases} C_1 & \text{if } g(\mathbf{x}) > 0 \\ C_2 & \text{otherwise.} \end{cases}$$

## Generalization to $K > 2$ Classes

- One way to handle multiple classes is to define  $K$  two-class problems, each separating one class from all other classes combined.
- Let  $g_i(\mathbf{x})$  denote  $\mathbf{w}_i^\top \mathbf{x} + \mathbf{w}_{i0}$ .
- An SVM  $g_i(\mathbf{x})$  is learned for each two-class problem.
- Classification rule during testing:

Classify to class  $j$  if  $j = \arg \max_{1 \leq k \leq K} g_k(\mathbf{x})$ .

## Relaxing the Constraints



- In practice, a separating hyperplane may not exist, possibly due to a **high noise level** which causes a large **overlap** of the classes.
- Even if a separating hyperplane exists, it is not always the best solution to the classification problem when there exist **outliers** in the data.
- A mislabeled example can become an outlier which affects the location of the separating hyperplane.

# Slack Variables

- A **soft-margin SVM** allows for the possibility of violating the inequality constraints

$$y^{(\ell)}(\mathbf{w}^\top \mathbf{x}^{(\ell)} + w_0) \geq 1$$

by introducing **slack variables**

$$\xi_\ell \geq 0, \ell = 1, \dots, N,$$


which store the deviation from the margin.

- **Relaxed separation constraints:**

$$y^{(\ell)}(\mathbf{w}^\top \mathbf{x}^{(\ell)} + w_0) \geq 1 - \xi_\ell.$$



## Relaxed Separation Constraints



- One data point is inside the margin on the wrong side of the hyperplane:

# Penalty

- By making  $\xi_\ell$  large enough, the constraint on  $(\mathbf{x}^{(\ell)}, y^{(\ell)})$  can always be met. In order not to obtain the trivial solution where all  $\xi_\ell$  take on large values, we should **penalize** them in the objective function.
- Four cases:
  - $\xi_\ell = 0$ : no problem with  $\mathbf{x}^{(\ell)}$  (**no penalty**)
  - $0 < \xi_\ell < 1$ :  $\mathbf{x}^{(\ell)}$  lies on the right side of the hyperplane but in the margin (**small penalty**)
  - $\xi_\ell = 1$ :  $\mathbf{x}^{(\ell)}$  lies at the hyperplane (**penalty between the previous and next cases**)
  - $\xi_\ell > 1$ :  $\mathbf{x}^{(\ell)}$  lies on the wrong side of the hyperplane (**large penalty**)
- Number of misclassifications:  $\#\{\xi_\ell > 1\}$
- Number of nonseparable instances:  $\#\{\xi_\ell > 0\}$
- **Soft error** as additional penalty term:

$$\sum_{\ell} \xi_\ell$$

# Primal and Dual Optimization Problems

- Primal optimization problem:

$$\begin{aligned} &\text{Minimize} && \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{\ell} \xi_{\ell} \\ &\text{subject to} && y^{(\ell)}(\mathbf{w}^{\top} \mathbf{x}^{(\ell)} + w_0) \geq 1 - \xi_{\ell}, \quad \forall \ell \\ &&& \xi_{\ell} \geq 0, \quad \forall \ell. \end{aligned}$$

- Dual optimization problem:

$$\begin{aligned} &\text{Maximize} && \sum_{\ell} \alpha_{\ell} - \frac{1}{2} \sum_{\ell} \sum_{\ell'} \alpha_{\ell} \alpha_{\ell'} y^{(\ell)} y^{(\ell')} (\mathbf{x}^{(\ell)})^{\top} \mathbf{x}^{(\ell')} \\ &\text{subject to} && \sum_{\ell} \alpha_{\ell} y^{(\ell)} = 0 \text{ and } 0 \leq \alpha_{\ell} \leq C, \quad \forall \ell. \end{aligned}$$

# Alternative View as Unconstrained Optimization Problems

- An alternative view is to formulate the learning of hard-margin and soft-margin SVMs as **unconstrained optimization** problems.
- Each optimization problem minimizes a **loss function** with a **regularizer**, i.e., a regularized loss function.
- Different loss functions are used for the hard-margin and soft-margin SVMs.

# Alternative Formulation for Hard-Margin SVM

- Objective function for minimization:

$$\sum_{\ell=1}^N E_{\infty}(y^{(\ell)} g(\mathbf{x}^{(\ell)}) - 1) + \lambda \|\mathbf{w}\|^2,$$

where

$$E_{\infty}(z) = \begin{cases} 0 & z \geq 0 \\ \infty & \text{otherwise.} \end{cases}$$

- Equivalent formulation that is more consistent with that of the soft-margin SVM:

$$\sum_{\ell=1}^N E_{\infty}(y^{(\ell)} g(\mathbf{x}^{(\ell)})) + \lambda \|\mathbf{w}\|^2,$$

where

$$E_{\infty}(z) = \begin{cases} 0 & z \geq 1 \\ \infty & \text{otherwise.} \end{cases}$$

# Alternative Formulation for Soft-Margin SVM


- Objective function for minimization:

$$\sum_{\ell=1}^N E_h(y^{(\ell)} g(\mathbf{x}^{(\ell)})) + \lambda \|\mathbf{w}\|^2,$$

where

$$E_h(z) = [1 - z]_+ = \max(0, 1 - z)$$

is called the **hinge loss**.



# Key Ideas of Kernel Methods

- Instead of defining a nonlinear model in the original (input) space, the problem is mapped to a new (feature) space by performing a **nonlinear transformation** using suitably chosen **basis functions**.
- A **linear** model is then applied in the new space.
- The basis functions are often defined **implicitly** via defining **kernel functions** directly.

# Basis Functions

- Basis functions:

$$\mathbf{z} = \phi(\mathbf{x}) \quad \text{where } z_j = \phi_j(\mathbf{x}), \quad j = 1, \dots, k.$$

- Discriminant function:

$$g(\mathbf{z}) = \mathbf{w}^\top \mathbf{z}$$

$$g(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}) = \sum_{j=1}^k w_j \phi_j(\mathbf{x}),$$

where we do not use a separate  $w_0$  but assume that  $z_1 = \phi_1(\mathbf{x}) \equiv 1$ .

- Usually,  $k \gg d, N$  (in fact  $k$  can even be infinite). The **dual form** is preferred because its complexity depends on  $N$  but that of the primal form depends on  $k$ .



# Kernel Functions

- Dual optimization problem:

$$\begin{aligned} &\text{Maximize} && \sum_{\ell} \alpha_{\ell} - \frac{1}{2} \sum_{\ell} \sum_{\ell'} \alpha_{\ell} \alpha_{\ell'} y^{(\ell)} y^{(\ell')} \phi(\mathbf{x}^{(\ell)})^{\top} \phi(\mathbf{x}^{(\ell')}) \\ &\text{subject to} && \sum_{\ell} \alpha_{\ell} y^{(\ell)} = 0 \text{ and } 0 \leq \alpha_{\ell} \leq C, \forall \ell, \end{aligned}$$

or

$$\begin{aligned} &\text{Maximize} && \sum_{\ell} \alpha_{\ell} - \frac{1}{2} \sum_{\ell} \sum_{\ell'} \alpha_{\ell} \alpha_{\ell'} y^{(\ell)} y^{(\ell')} K(\mathbf{x}^{(\ell)}, \mathbf{x}^{(\ell')}) \\ &\text{subject to} && \sum_{\ell} \alpha_{\ell} y^{(\ell)} = 0 \text{ and } 0 \leq \alpha_{\ell} \leq C, \forall \ell, \end{aligned}$$

where  $K(\mathbf{x}^{(\ell)}, \mathbf{x}^{(\ell')}) \equiv \phi(\mathbf{x}^{(\ell)})^{\top} \phi(\mathbf{x}^{(\ell')})$  is a **kernel function** defined directly on the inputs  $\mathbf{x}^{(\ell)}$  and  $\mathbf{x}^{(\ell')}$ .

## Some Common Kernel Functions: Polynomial Kernel

- Polynomial kernel:

$$K(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}' + 1)^q,$$

where  $q$  is the degree.

- E.g., when  $q = 2$  and  $d = 2$ ,

$$\begin{aligned} K(\mathbf{x}, \mathbf{x}') &= (\mathbf{x}^\top \mathbf{x}' + 1)^2 \\ &= (x_1 x'_1 + x_2 x'_2 + 1)^2 \\ &= 1 + 2x_1 x'_1 + 2x_2 x'_2 + 2x_1 x_2 x'_1 x'_2 + (x_1)^2 (x'_1)^2 + (x_2)^2 (x'_2)^2, \end{aligned}$$

which corresponds to the inner product of the basis function

$$\phi(\mathbf{x}) = (1, \sqrt{2}x_1, \sqrt{2}x_2, \sqrt{2}x_1x_2, (x_1)^2, (x_2)^2)^\top.$$

## Some Common Kernel Functions: RBF Kernel

- Radial basis function (RBF) kernel:

$$K(\mathbf{x}, \mathbf{x}') = \exp \left( -\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2s^2} \right).$$

- It can be generalized to

$$K(\mathbf{x}, \mathbf{x}') = \exp \left( -\frac{\mathcal{D}(\mathbf{x}, \mathbf{x}')}{2s^2} \right),$$

where  $\mathcal{D}(\cdot, \cdot)$  is some distance function.

## To Learn More...

- One-class support vector machines